

```
In [1]: import tensorflow as tf
```

Loading - "MNIST Data Set"

Contains Training samples = 60,000 and Testing samples = 10,000

from tensorflow

```
In [2]: mnist = tf.keras.datasets.mnist #handwritten characters
```

Divide into train and test datasets

```
In [3]: #unpacking the dataset into train and test datasets  
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11493376/11490434 [=====] - 7s 1us/step


11501568/11490434 [=====] - 7s 1us/step

```
In [4]: x_train.shape
```


```
Out[4]: (60000, 28, 28)
```

```
In [6]: import matplotlib.pyplot as plt
```

```
In [7]: #checking how data looks like  
plt.imshow(x_train[0])  
plt.show() #executing the graph  
#we donno whether it's color image or binary images  
#so inorder to plot it, change the config  
plt.imshow(x_train[0], cmap = plt.cm.binary)
```



```
Out[7]: <matplotlib.image.AxesImage at 0x1b8437be520>
```



checking the values of each pixel

Before Normalization

In [8]: `print(x_train[0]) #before normalization`

```
[
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  3  18  18  18 126 136
175 26 166 255 247 127 0 0 0 0
  0  0  0  0  0  0  0  0  30 36 94 154 170 253 253 253 253
225 172 253 242 195 64 0 0 0 0
  0  0  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251
93 82 82 56 39 0 0 0 0 0
  0  0  0  0  0  0  0 18 219 253 253 253 253 198 182 247 241
0 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 80 156 107 253 253 205 11 0 43 154
0 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 14 1 154 253 90 0 0 0 0
0 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 139 253 190 2 0 0 0
0 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 11 190 253 70 0 0 0
0 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 0 35 241 225 160 108 1
0 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 0 0 81 240 253 253 119
25 0 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 0 0 45 186 253 253
150 27 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 0 0 0 16 93 252
253 187 0 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 0 0 0 0 249
253 249 64 0 0 0 0 0 0 0
  0  0  0  0  0  0  0 0 0 0 0 0 0 0 46 130 183 253
253 207 2 0 0 0 0 0 0 0]
```

Images are in Gray level (1 pixel == 0 to 255), not (RGB)

[illegible]

4/10

5/10

```
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.]
```

```
In [13]: print(y_train[0]) #labels
```

5

Resizing image to make it suitable for applying Convolution operation

```
In [15]: import numpy as np
IMG_SIZE = 28
x_trainr = np.array(x_train).reshape(-1, IMG_SIZE, IMG_SIZE,1) ### increasing 1 dime
x_testr = np.array(x_test).reshape(-1, IMG_SIZE, IMG_SIZE,1) ### increasing 1 dimens
print("Training Samples dimension", x_trainr.shape)
print("Testing Samples dimensions",x_testr.shape)
```

Training Samples dimension (60000, 28, 28, 1)

Testing Samples dimensions (10000, 28, 28, 1)

Creating a Deep Neural Network

Training on 60,000 samples of MNIST handwritten dataset

```
In [17]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, Max
```

```
In [20]: ### creating a neural network
model = Sequential()

#### First Layer
model.add(Conv2D(64, (3,3), input_shape = x_trainr.shape[1:])) ### mention input lay
model.add(Activation('relu')) ## Activation function, to make it non-linear
model.add(MaxPooling2D(pool_size = (2,2))) ## Maxpooling, Single max value of 2x2 ma

#### Second Layer
model.add(Conv2D(64, (3,3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))

#### Third Layer
model.add(Conv2D(64, (3,3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))

#### Fully Connected Layer #1
model.add(Flatten()) ### before using fully connected, it need to be flatten (2D to
model.add(Dense(64)) # neural network
model.add(Activation('relu'))

#### Fully Connected Layer #2
model.add(Dense(32))
model.add(Activation('relu'))

#### Last Fully Connected Layer, output must be equal to number of classes, 10(0-9)
model.add(Dense(10)) ## Last dense layer must be equal to 10
model.add(Activation('softmax')) ## changed to sotmax (class probabilites)
```

In [22]:

```
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 26, 26, 64)	640
activation_2 (Activation)	(None, 26, 26, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 64)	0
conv2d_3 (Conv2D)	(None, 11, 11, 64)	36928
activation_3 (Activation)	(None, 11, 11, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_4 (Conv2D)	(None, 3, 3, 64)	36928
activation_4 (Activation)	(None, 3, 3, 64)	0
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 64)	0
flatten (Flatten)	(None, 64)	0
dense (Dense)	(None, 64)	4160
activation_5 (Activation)	(None, 64)	0
dense_1 (Dense)	(None, 32)	2080
activation_6 (Activation)	(None, 32)	0
dense_2 (Dense)	(None, 10)	330
activation_7 (Activation)	(None, 10)	0
=====		
Total params: 81,066		
Trainable params: 81,066		
Non-trainable params: 0		

In [23]:

```
print("Total Training Samples = ", len(x_train))
```

Total Training Samples = 60000

In [24]:

```
model.compile(loss = "sparse_categorical_crossentropy", optimizer = "adam", metrics=
```

In [25]:

```
model.fit(x_train, y_train, epochs=5, validation_split = 0.3) ## Training my model
```

Epoch 1/5

1313/1313 [=====] - 26s 19ms/step - loss: 0.3518 - accuracy: 0.8880 - val_loss: 0.1345 - val_accuracy: 0.9591

Epoch 2/5

```

1313/1313 [=====] - 24s 19ms/step - loss: 0.1119 - accurac
y: 0.9655 - val_loss: 0.1261 - val_accuracy: 0.9623
Epoch 3/5
1313/1313 [=====] - 26s 20ms/step - loss: 0.0807 - accurac
y: 0.9749 - val_loss: 0.0788 - val_accuracy: 0.9749
Epoch 4/5
1313/1313 [=====] - 26s 20ms/step - loss: 0.0622 - accurac
y: 0.9800 - val_loss: 0.0675 - val_accuracy: 0.9793
Epoch 5/5
1313/1313 [=====] - 26s 20ms/step - loss: 0.0503 - accurac
y: 0.9837 - val_loss: 0.0705 - val_accuracy: 0.9788
Out[25]: <keras.callbacks.History at 0x1b84bc786d0>

```

In [26]:

```

### Evaluating on testing dataset MNIST
test_loss, test_acc = model.evaluate(x_testr, y_test)
print("Test loss on 10,000 test samples ",test_loss)
print("Validation Accuracy on 10,000 test samples ",test_acc)

```

```

313/313 [=====] - 2s 5ms/step - loss: 0.0673 - accuracy: 0.
9796
Test loss on 10,000 test samples  0.06731598824262619
Validation Accuracy on 10,000 test samples  0.9796000123023987

```

In [27]:

```

predictions = model.predict([x_testr])

```

In [28]:

```

print(predictions)

```

```

[[1.0075364e-08 4.1938529e-07 7.2846314e-07 ... 9.9998963e-01
 3.7899053e-08 8.1203061e-06]
 [2.2520733e-03 8.1896333e-06 9.9697220e-01 ... 1.1457161e-05
 7.6291493e-05 8.9106429e-08]
 [2.7473826e-07 9.9988329e-01 1.0038377e-06 ... 1.7928201e-05
 5.5844421e-06 4.0396847e-07]
 ...
 [1.0216693e-10 4.6228013e-07 3.9864737e-10 ... 1.5555089e-07
 1.0085337e-06 1.5766143e-06]
 [1.5050702e-04 2.5392904e-07 8.8294797e-08 ... 1.0789056e-06
 1.2867320e-03 1.0875900e-05]
 [8.5594296e-04 2.5891336e-07 4.8824097e-04 ... 8.7840294e-07
 1.0400967e-04 2.8484084e-05]]

```

In [29]:

```

print(np.argmax(predictions[0]))

```

7


In [30]:

```

plt.imshow(x_test[0])

```

Out[30]: <matplotlib.image.AxesImage at 0x1b84b671f70>




```
In [31]: print(np.argmax(predictions[128]))
```

8

```
In [32]: plt.imshow(x_test[128])
```

```
Out[32]: <matplotlib.image.AxesImage at 0x1b83e991ee0>
```




```
In [33]: import cv2 #importing own image to check the predictions
```

```
In [36]: img = cv2.imread('eight.png')
```

```
In [37]: plt.imshow(img)
```

```
Out[37]: <matplotlib.image.AxesImage at 0x1b84124ea00>
```



```
In [39]: img.shape
```

```
Out[39]: (296, 296, 3)
```

```
In [40]: gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
gray.shape
```

```
Out[40]: (296, 296)
```

```
In [48]: resized = cv2.resize(gray, (28,28), interpolation = cv2.INTER_AREA)
```

```
In [49]: resized.shape
```

```
Out[49]: (28, 28)
```

```
In [50]: newimg = tf.keras.utils.normalize (resized, axis=1)
```

```
In [51]: newimg = np.array(newimg).reshape(-1, IMG_SIZE, IMG_SIZE, 1)
```

```
In [52]: newimg.shape
```

```
Out[52]: (1, 28, 28, 1)
```

```
In [53]: predictions = model.predict(newimg)
print(np.argmax(predictions))
```

8